



Vietnamese - German University



VIETNAMESE - GERMAN UNIVERSITY
FRANKFURT UNIVERSITY OF APPLIED SCIENCES

FACULTY OF ENGINEERING
COMPUTER SCIENCE DEPARTMENT

BACHELOR THESIS

Evaluation And Implementation of Modern Virtual Private Network Protocol for The GNRC Network Stack

By

Le Hoang Dang Nguyen

Matriculation number: 17028

Assessor:

1. Prof. Dr. Oliver Hahm
2. Dr. Tran Thi Thu Huong

Binh Duong, 2024

This page intentionally left blank.

Declaration

I hereby declare that this thesis is a product of my own work, unless otherwise referenced. I also declare that all opinions, results, conclusions and recommendations are my own and may not represent the policies or opinions of Vietnamese – German University.

Le Hoang Dang Nguyen

Acknowledgements

I want to thank you.

Abstract

Abstract go here

Contents

1	Introduction	9
1.1	Motivation	9
1.2	Organization	10
2	Wireless Embedded Internet	12
2.1	Overview	12
2.2	The 6LoWPAN Architecture	14
2.3	6LoWPAN Protocol Stack	17
2.4	IEEE 802.15.4	18
3	RIOT-OS and The GNRC Network Stack	20
3.1	RIOT Operating System	20
3.2	GNRC	23
3.2.1	Overview & Architecture	23
3.2.2	The Packet Buffer - pktbuf	25
3.2.3	GNRC's Module Registry - netreg	27
3.2.4	Network Interfaces	28
4	Concepts of Wireguard	29

4.1	Protocol & Cryptography	29
4.1.1	Overview	29
4.1.2	Cryptokey Routing	33
4.2	Noise	34
4.2.1	IKpsk2 Handshake Pattern	40
4.3	Wireguard Messages	45
4.4	Timers	49
5	Wireguard and the IoT	50
5.1	Cryptography Resilience	50
5.2	Virtuous Silence	50
5.3	Stateful firewall	50
5.4	Cookies and Denial of Services Attack	50
5.5	Roaming & NDP	50
5.6	Message Lengths & Fragmentation	50
6	Related Work	51
7	Design and Implementation of Wireguard for GNRC	52
8	Testing	53
8.1	Methodology	53
8.2	Environment	53
8.3	Scenario 1	53
8.4	Scenario 2	53
8.5	Scenario 3	53

8.6	Scenario 4	53
8.7	Memory Usage	53
9	Conclusion and Future Work	54

Chapter 1

Introduction

1.1 Motivation

As VPN is mentioned in many papers as not suitable for constrained small IoT Devices, I start to question, is it really the true for in this modern area. As a modern network protocols called wireguard show up, with an aim to replace the old, sound but complex protocols like IPSec, we here consider adapting the wireguard protocols onto this constrained IoT environment, to see whether this VPN protocol is a good fit for the IoT world or not.

WireGuard [23] is a new VPN protocol that fits the role of this new pipe and it looks quite promising. Note that WireGuard was originally presented at NDSS 2017 [15], but while the main concepts still apply, the protocol has slightly evolved in an incompatible way. The latest version is described in the WireGuard whitepaper [18]. WireGuard is designed using modern cryptography, aims for high performance and reduces the attack surface as a simple protocol. Unlike other

protocols, no form of negotiation over cryptographic parameters is possible. Instead, it uses the following constructions and algorithms:

Noise protocol framework A collection of cryptographic handshake patterns which provide building blocks to construct new secure protocols with authenticated key agreement [47]. ChaCha20-Poly1305 The ChaCha20 stream cipher and Poly1305 authenticator, used in an Authenticated Encryption with Additional Data (AEAD) construction. This is also used in modern security protocols such as TLS 1.3 [45]. It provides authenticity and confidentiality of transported data.

X25519 An elliptic-curve-Diffie-Hellman (ECDH) function [4]. Compared to other functions, it has a rather small key size and simpler requirements regarding key validation. It is used in the key agreement protocol.

HKDF The HMAC-based Key Derivation Function (HKDF) is a construction to derive one or more keys from an initial secret [40]. It is used to link all pieces of the handshake state to each other, including keys and protocol messages. It also ensures that the original key material that is involved in calculations cannot be recovered.

BLAKE2 A fast cryptographic hash function, BLAKE2s [50], is used by the HKDF and as a message authentication code (MAC).

1.2 Organization

Chapter 2 discusses about the background knowledge of IoT, RIOT Operating system, and the 6LoWPAN network stack.

Chapter 3 is about the underlying network stack that powers RIOT - GNRC.

Chapter 4 is a thorough explanation of the wireguard protocol.

Chapter 5 gives an evaluation of wireguard security, and how it help the IoT world.

Chapter 6 gives some related works of other compressed VPN protocols, and other evaluation of wireguard for constrained devices.

Chapter 7 discusses the design and implementation of wireguard.

Chapter 8 explains the testing of the wireguard.

Chapter 9: conclusion and future works of this thesis.

Chapter 2

Wireless Embedded Internet

This chapter provide the definition and overview of the wireless embedded internet. It covers It covers architectures specifically developed for IoT applications, highlights the 6LoWPAN protocol stack that allows the integration of the internet protocol (IP) stack onto this type of network, and reviews 802.15.4, a common protocol in embedded context.

2.1 Overview

The Internet of Things (IoT) comprises of all embedded devices and networks that are natively IP-enabled and connected to the Internet, such as sensors, machines, and RFID readers, alongside the services monitoring and controlling these devices [SB09]. A subset of IoT, the Wireless Embedded Internet consists of low-power, resource-limited wireless devices connected through standards like IEEE 802.15.4. Integrating standard Internet protocols with such networks presents sev-

eral challenges:

Power and duty-cycle: The IP-enabled devices should always be connecting, contradicting the low-duty-cycle nature of the battery-powered wireless devices.

Multicast: Wireless embedded radio technologies like IEEE 802.15.4, do not generally support multicast, and flooding wastes power and bandwidth in such network. However, Multicast is an important operation for many IPv6 features such as address auto-configuration [NJT07].

Limited bandwidth and frame sizes: Bandwidth and frame size inside a low-power wireless embedded radio network are limited, with only 20-250 kbits/s and 40-200 bytes correspondingly. For example, the frame size for IEEE 802.15.4 standard has a 127-byte frame size, with layer-2 payload sizes as low as 72 bytes. The minimum frame size for standard IPv6 is 1280 bytes [DH17], hence fragmentation is required.

Reliability: Standard Internet Protocols are not optimized for low-power wireless networks. For example, TCP is not able to differentiate between packets dropped by congestion or lost on wireless links. Node failure, energy exhaustion, and sleep duty cycles can also incur unreliability in wireless embedded networks.

To tackle these issues, 6LoWPAN [Mon+07] was developed, enabling IPv6 and its related protocols to function effectively in wireless

embedded networks. IPv6's simple header structure and hierarchical addressing make it ideal for use in these constrained environments.

2.2 The 6LoWPAN Architecture

According to Zach Shelby and Carsten Bormann, the Wireless Embedded Internet is formed by connecting islands of wireless embedded devices, with each island functioning as a stub network within the Internet [SB09, p. 13]. A stub network is one where IP packets are either sent to or received from, but it does not serve as a transit point for other networks. The 6LoWPAN architecture is illustrated in Figure 2.1. In this context, the 6LoWPAN architecture consists of low-power wireless area networks (LoWPANs), which operate as IPv6 stub networks. Each LoWPAN is a set of 6LoWPAN nodes, sharing a common IPv6 address prefix (the first 64 bits of an IPv6 address), with the interconnection between the LoWPANs achieved through the edge router. There are 3 different kinds of LoWPANs:

- An **Ad hoc LoWPAN** operates independently without the connection to the internet.
- A **Simple LoWPAN** connects to another IP network via an edger router.
- An **External LoWPAN** comprises LoWPANS of multiple edge routers along with a backbone link connecting them.




Figure 2.1: The 6LoWPAN architecture, see [SB09, p. 14]

LoWPANs are connected to other IP networks via edge routers, as illustrated in Figure 2.1. The edge router plays a key role by routing traffic to and from the LoWPAN, managing 6LoWPAN compression, handling Neighbor Discovery within the network, and other network management features. Each node in a LoWPAN could either be a host, an edge router, or a node routing between other nodes. The shared common IPv6 prefix within the LoWPAN is advertised by edge routers or is pre-configured on each node. An edge router keeps a list of registered nodes that are accessible through its network interface inside the LoWPAN.

To enter a 6LoWPAN, a node sends a Router Solicitation message to obtain the IPv6 prefix unless it has been statically configured. Upon receiving the prefix, the node generates a unique global IPv6 address and registers this address with the edge router of the LoWPAN. This allows the edge router to make informed routing decisions for traffic entering and exiting the LoWPAN, as well as to facilitate neighbor discovery within the 6LoWPAN. The edge router must update the list of registered nodes regularly, as addresses expire after a configurable period. A longer expiration time helps reduce a node's power consumption, while a shorter expiration time accommodates rapidly changing network structures. These processes are defined in the dedicated neighbor discovery protocol for 6LoWPAN [Bor+12]. LoWPAN nodes can travel freely within and among multiple LoWPAN networks, and they may participate in several LoWPANs simultaneously. Communication between a LoWPAN node and an external IP node occurs in an end-to-end manner, similar to interactions between standard IP nodes.

In an extended LoWPAN, multiple edge routers are integrated into the same LoWPAN, sharing the same IPv6 prefix. These edge routers are interconnected through a common backbone link. When a node moves between edge routers, it must register with the edge router it can access, but it retains its IPv6 address. Communication between edge routers regarding neighbor discovery is handled over the backbone link, which reduces messaging overhead. This extended LoW-

PAN architecture allows a single LoWPAN to cover larger areas.

An ad-hoc LoWPAN works in the same manner as a simple LoWPAN, but without the link to other IP networks. Instead of an edge router, a node will act as a simple edge router, handle unique local address generation [HH05] and provide the neighbor discovery registration feature to other nodes.

2.3 6LoWPAN Protocol Stack

Figure 2.2 depicts the IPv6 protocol stack with 6LoWPAN in comparison to a standard IP protocol stack and the corresponding five layers of the Internet Model. This Internet Model connects a wide range of link-layer technologies with various transport and application protocols.




Figure 2.2: IP and 6LoWPAN stack, see [SB09, p. 16]

The IPv6 protocol stack with 6LoWPAN (sometimes called the 6LoWPAN protocol stack) is nearly identical to a conventional IP stack, with a few notable differences. Primarily, 6LoWPAN only supports IPv6, for which a small adaptation layer—the LoWPAN adapta-

tion layer—has been defined to optimize IPv6 for IEEE 802.15.4 and similar link layers, as detailed in [Mon+07]. In practice, many implementations of the 6LoWPAN stack in embedded devices combine the LoWPAN adaptation layer with IPv6, allowing them to be displayed together as part of the network layer.

The User Datagram Protocol (UDP) [Pos80] is the most commonly used transport protocol with 6LoWPAN, and it can be compressed using the LoWPAN format. The Transmission Control Protocol (TCP) is less frequently used due to concerns about performance, efficiency, and complexity, although there have been recent efforts on guidance to use and implement TCP for the IoT [GCS21]. The Internet Control Message Protocol version 6 (ICMPv6) [GC06] is utilized for control messaging, including functions like ICMP echo requests, destination unreachable messages, and Neighbor Discovery. While many application protocols are application-specific and in binary format, there is a growing availability of more standardized application protocols.

2.4 IEEE 802.15.4

Established by the IEEE, the IEEE 802.15.4 standards define low-power wireless radio techniques and specify the physical and media access control layers that serve as the foundation for 6LoWPAN. The IEEE 802.15.4-2011 version of the standards includes features such as access control via CSMA/CA, optional acknowledgments for re-

transmission of corrupted data, and 128-bit AES encryption at the link layer. It offers addressing modes that utilize both 64-bit and 16-bit addresses with unicast and broadcast capabilities. The payload of a physical frame can reach up to 127 bytes, with 72 to 116 bytes of the usable payload after link-layer framing, depending on different addressing and security options [SB09, Appendix B.1].

Star and point-to-point network topologies are supported by IEEE 802.15.4. The MAC layer can operate with CSMA/CA in the beaconless mode. In beacon-enabled mode, TDMA/TISCH for media access is utilized. The number of nodes, the length of the transmitted messages, and the level of radio interference within the ISM band significantly influenced the average packet loss in IEEE 802.15.4 [Shu+07]. While the use of acknowledgments at the link layer enhances reliability, it complicates the estimation of packet round-trip times.

Chapter 3

RIOT-OS and The GNRC Network Stack

3.1 RIOT Operating System

RIOT, the friendly operating system for the Internet of Things, is a real-time operating system, specifically designed for low-end IoT devices with a minimal memory in order of $\approx 10\text{K}$ Byte [Bac+18]. It can run on devices with neither memory management unit (MMU) nor memory protection unit (MPU).

Under the distribution of LGPLv2.1 License, RIOT is free and open-source software, meaning it can be used and distributed by anyone. Furthermore, this license allows the linkage of RIOT with proprietary software and supports the ability to be customized by the end users.

The design objectives of RIOT focus on several key areas: optimizing resource usage such as RAM, ROM, and power consumption;

supporting a broad spectrum of configurations, from 8-bit to 32-bit MCUs, and accommodating various boards and use cases; reducing code duplication across different setups; ensuring most of the code is portable across supported hardware; offering user-friendly software platform; and enabling real-time capabilities. To realize these goals, one of the principles that the RIOT follows is modularity.

RIOT is organized into software modules that are combined at compile time, centered around a kernel offering minimal functionality. This modular approach allows the system to be built in a way that includes only the necessary modules for a given use case. As a result, both memory usage and system complexity are kept to a minimum in practical deployments. The code structure of RIOT is illustrated in Figure 3.1:

- **core** provides the kernels and basic data structures like linked lists, LIFOs, and ringbuffers.
- Four parts of hardware abstractions:
 1. **cpu** implements functionalities of microcontroller.
 2. **boards** selects, maps and configures the used CPU and drivers.
 3. **drivers** implement the device drivers.
 4. **periph** provides unified access to microcontroller peripherals and is used by device drivers.
- **sys** implements libraries beyond kernel features, such as cryptog-

raphy, networking, and file system.

- **pkg** contains third-party libraries which do not exist within the main code repository.

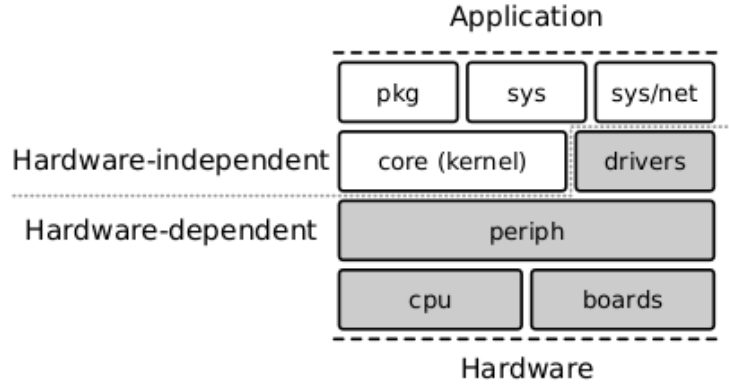


Figure 3.1: Structure elements of RIOT, see [Bac+18, p. 3]

Multi-threading is a builtin feature for RIOT to offer several benefits: (a) clear logical separation between different tasks, (b) straightforward task prioritization, and (c) easier integration of external code [Bac+18, p. 4]. Various synchronization primitives, such as mutex, semaphore, and message passing (**msg**) are provided by RIOT kernel. Multi-threading can also be optional in the case of extremely low-memory usage of the application.

RIOT's kernel employs a scheduler that uses fixed priorities and preemption with $O(1)$ operations, enabling soft real-time capabilities. Specifically, the time required to interrupt and switch between threads is bounded by a small upper limit, as operations like context saving, selecting the next thread, and context restoring are deterministic. The system follows a class-based run-to-completion scheduling policy,

where the highest-priority active thread is executed and can only be interrupted by interrupt service routines (ISRs). This scheduler allows RIOT to effectively prioritize tasks, ensuring that high-priority events can preempt lower-priority tasks as needed.

RIOT's scheduler operates in a tickless manner, meaning it does not rely on CPU time slices or periodic system timer ticks. As a result, the system remains in a low-power state unless an actual event occurs, such as an interrupt triggered by hardware. Wake-up events can be initiated by a transceiver receiving a packet, timers expiring, buttons being pressed, or similar activities. When no threads are in a running state and no interrupts are pending, the system automatically switches to the idle thread, which has the lowest priority. The idle thread, in turn, transitions the system into the most energy-efficient mode available thereby reducing energy consumption.

3.2 GNRC

3.2.1 Overview & Architecture

The GNRC (short for generic) network stack is the default network stack for RIOT [Len16]. It was designed as a replacement for the previous monolithic and unmaintainable network stack of RIOT [Bru16]. The old stack was lack of unified interfaces, making it difficult to achieve modularity, testability, and extensibility. Furthermore, as most of the networks were using their own buffers, the memory foot-

print was increased, creating a need for significant data copying between layers, thus leading to a reduction in overall performance. These limitations drove the creation of the "gnrc" network stack to replace the previous one.

The design objectives of the generic network stack encompass a low memory footprint, full-featured protocols, modular architecture, support for multiple network interfaces, parallel data handling and customization during compilation time.




Figure 3.2: GNRC Network Stack

Figure 3.2 illustrates the architecture of the GNRC network stack. Using multi-threading features and RIOT's thread-targeted IPC, every networking module runs in its own thread, communicating with other modules via the message queue. The messages passing between the network threads follow a clearly defined format called "*netapi*" in GNRC. A network registry called "*netreg*" is used for a network thread to search for a module that is interested in receiving the next packet, using the packet's type. Modules that want to receive cer-

tain types can register with this registry. While traversing through the stack, a packet is stored in the centralized GNRC packet buffer “*pktbuf*”. With this architecture, GNRC attains the required modularity and simplifies testing. It also offers an easy way to prioritize different components of the stack, by assigning different priorities to the threads running the protocols and allowing the operating system’s process scheduler to handle these priorities. The trade-off of this design is a performance penalty compared to straightforward function calls, as IPC involves context switches and context saves. However, this overhead is shown to be only one order of magnitude slower than a direct function call on real IoT hardware [Pet+15, section 4.1].

3.2.2 The Packet Buffer - *pktbuf*

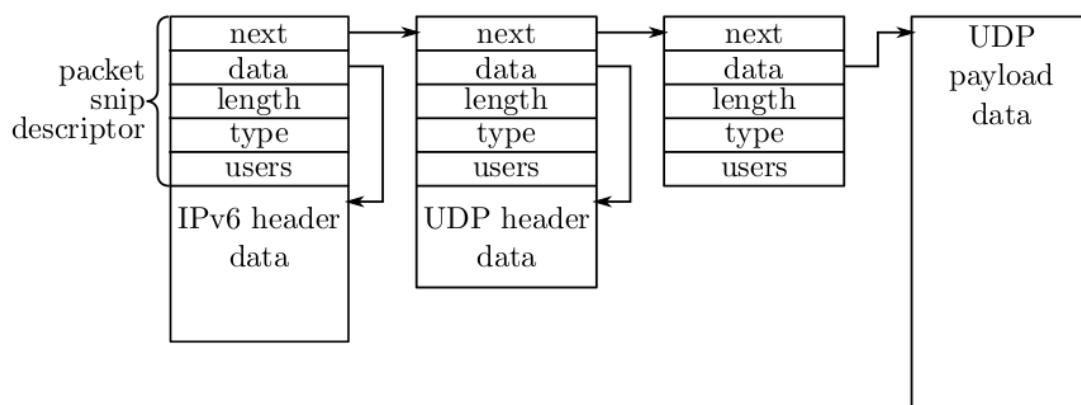


Figure 3.3: An example of a packet in transmission in the GNRC packet buffer, see [Len16, p. 22]

Packets in GNRC are stored in a packet buffer called *pktbuf*, consisting of variable-length segments known as “packet snips” (see Figure 3.3). These snips are typically used to distinguish between different

headers and payloads. The packet data, along with the structures describing the packet snips, is stored in arrays allocated in static memory, although the API can theoretically support dynamic allocation on the memory heap. Packet snips can be labeled with a type to indicate the content of the packet. To create a packet snip in `pktbuf`, a user can use the function `gnrc_pktbuf_add()`, which takes all user-controllable fields of a packet snip (`next`, `data`, `length`, and `type`) as arguments and returns the new snip.

With minimal packet duplication across the stack, data is being copied or moved as infrequently as possible from the network interface to the application, ideally only once at each end-point of the stack and vice versa. This approach allows the structures that describe the packet (referred to as the “packet snip descriptor” in Figure 3.3) to be stored independently of the actual data, enabling the marking of a packet header without requiring to move data around.

To ensure safety in concurrency, a reference counter (`user`) is maintained in a packet snip to keep track of the using threads. The APIs that handle the increment and the decrement of the counter are `gnrc_pktbuf_hold()` and `gnrc_pktbuf_release()`. The packet is removed from the packet buffer when the counter hits 0.

When a thread starts to write to a packet and the reference counter is greater than one, the packet snip and all its next pointer will be cloned using a copy-on-write approach, resulting in the decrement of the original snip’s reference counter by one. For minimal duplication,

the pointers reaching the current snip will be copied, creating a tree structure for the packet within the packet buffer, hence reversing the order of packet snips for a received packet. The packet snips now starts with the payload and end with the header of the lowest layer, whilst a packet in transmission retains the order in which it will be sent, reducing the number of pointers needed for the tree structure. Nevertheless, it's crucial to note that data will be typically kept in the order it was received by the packets in reception. The only thing to be reversed is the packet descriptor list that marks the data. When a packet (i.e., the first snip in a packet snip list) is released by a thread, its own snips will be removed from the packet buffer, with any of its copy snips still persist. The implementation of this feature is implemented in the `gnrc_pktbuf_start_write()` function.

This certain way of storing packets brings the merits of parallel data handling and low-memory footprint to GNRC, while avoiding duplication across the network stack.

3.2.3 GNRC's Module Registry - `netreg`

The *netreg* API provides a central directory to the generic network stack. When the thread of a network module is created, it also registers in *netreg* using its thread PID and "NETTYPE" - the kind of information that it is interested in. For example, an IPv6 module registers with its PID and type "NETTYPE_IPV6". If the UDP module wants to dispatch the packet down the stack, *netreg* would be used

to search for threads registered with type “NETTYPE_IPV6”. Using *netapi*, the module then send a pointer to the pktbuf-allocated packet to every interested thread.

3.2.4 Network Interfaces

In GNRC, network interfaces run in their own threads and communicate with the network layer via *netapi*. Within a network interface thread, MAC protocol are implemented and interact directly with network device drivers through the *netdev* API. This direct access are required since some link-layer protocols, such as TDMA-based MAC, need minimal delay when accessing the device. To reduce the overhead of dealing with L2 headers, the header formats of each network interface’s L2 protocol are converted into netif *netif* header - a general interface header format. Source and destination L2 addresses, header’s length, along side with additional link metrics utilized for routing protocol, such LQI and RSSI are included in this *netif* header. A unified conversion API for popular L2 protocols ensures that no redundant porting efforts are necessary.

Chapter 4

Concepts of Wireguard

This chapter provide the core concepts of Wireguard protocol, especially in its handshake and timer state machine. Section 4.1 gives a high-level overview of the traits of the protocols, and the cryptography primitives that they build upon. Section 4.2 details the foundational handshake protocol that Wireguard’s key exchange protocol builds upon. Section 4.3 defined all messages types of Wireguard and section 4.4 describes the timer state machine.

4.1 Protocol & Cryptography

4.1.1 Overview

Wireguard works as an encrypted IP network tunnel, resides at the layer 3 of the OSI layer, and uses UDP as its transport protocol. The establishment of a secure session before any transmission of data is via 1-RTT key-exchange handshake protocols. This handshake protocol is designed based on a Trevor Perin’s noise handshake pattern [Per].

After the handshake, the transported IP payload are protected using ChaCha20-Poly1305 Authenticated Encryption with Associated Data (AEAD) [NL18].

In Wireguard, the endpoints of communication have no role of server and client, but works in a peer-to-peer style. At any point in time, a peer can have a role of an **initiator**, start to send a handshake initiation message to create a secure channel with a **responder**. If the secure channel is not active for some amount of time, the initiator peer can change to a responder in the event of a the previous responder tries to initiate a new handshake, leading to the establishment of a new session between 2 peers. Figure 4.1 illustrates this handshake flow with periodic key rotation.

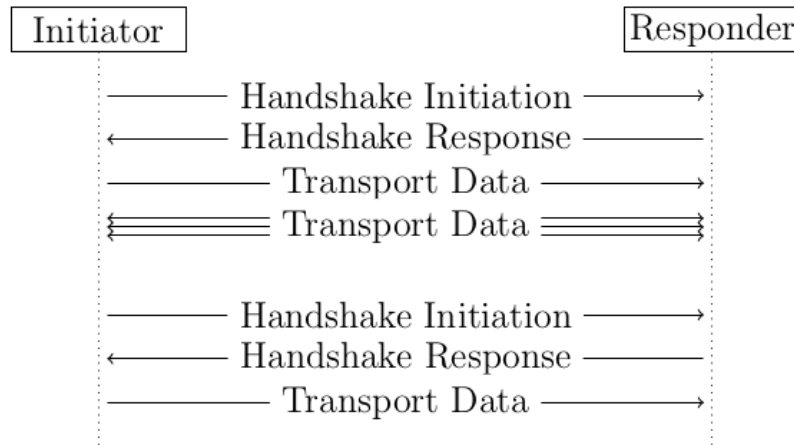


Figure 4.1: After the handshake, both peers can send data towards each other. After a duration, rekeying occurs to create a new session. See [Wu19, p. 7]

A static 32-byte Curve25519 [Ber06] public key are used to identified an endpoint inside the tunnel. Without the proof of sender's knowledge on this public key, an endpoint will never respond. Thus, adversaries can't probe any system protected by Wireguard, using

port scanning if they do not know the long-term static public key.

The Wireguard handshake protocol can be considered as a 1.5 round-trip time (1.5-RTT) handshake [Wu19]. Upon the reception of the handshake response after sending a handshake initiation, the initiator can promptly begin delivering encrypted payload. To safeguard against replay attacks, the responder must wait to send encrypted data messages until it has received an encrypted data message from the initiator, which serves as an acknowledgment of the handshake response.

During a handshake, new ephemeral key pairs will be generated by both parties, with the private key being wiped after the handshake, assuring the forward secrecy of the session. Each session is bound by fixed lifetime and the number of data messages that can be sent. When either of the limits is reached, the derivation of new session keys are required if the peers want to extend the communication. The impact of compromised session key is also mitigated by frequent periodic key rotation.

There is no direct shutdown signal within the protocol, only session removal after a fixed duration. Even if one side has terminated its VPN tunnel, the peer would remain unaware and keep forwarding data.

Wireguard also provides an optional pre-shared symmetric key (PSK) mode, where any pairs of peer pre-share a 256-bit symmetric encryption key between themselves, to assure the post-quantum security as long as the PSK never leaks out. This protection is against the idea

that if adversaries may be recording all the traffic for a long time, until the quantum computer exists. With PSK, despite the fact they maybe break all Curve25519-encrypted traffic, but not the ones that include a PSK.

Another adversary model that Wireguard mitigate is the denial of service through CPU exhaustion denial of services attack [Hen11, p. 268]. This mechanism is achieved via an encrypted cookie, when a peer is currently under load. This will be explained more in section 5.4.

At last, Wireguard is cryptographically opinionated. The protocol intentionally does not provide any form of flexibility when it comes to choosing the cryptography suite, hence a negation for cryptography algorithms does not exist. Specifically, Wireguard uses the following modern cryptography constructions:

Noise protocol framework A set of cryptographic handshake patterns that serve as building blocks for creating new secure protocols with authenticated key agreement.

Elliptic Curve Diffie-Hellman (ECDH) A Curve25519-based key-agreement protocol with small key size and simple requirements regarding key validation.

ChaCha20-Poly1305 a modern AEAD combining ChaCha20 [Ber08] stream cipher and Poly1305 [Ber05] authenticator to achieve authenticity and confidentiality.

XChaCha-Poly1305 an extended version of chacha20-poly1305 that supports 192-bit nonce. This is used for cookie encryption [Arc20].

HKDF The HMAC-based Key Derivation Function [KE10] to derive the encryption and decryption keys from the handshake state, keys and protocol messages.

BLAKE2 A fast and cryptographic hash function for message authentication code and is used by the HKDF [SA15].

4.1.2 Cryptokey Routing

The binding between peers and the allowed source IP addresses is the fundamental principle to build a secure VPN. As mentioned above, in Wireguard, peer must be identified by the 32-byte Curve25519 public key, allowing an association between the public key and the set of allowed IP addresses on the peer. This simple mapping is the concept behind the cryptokey routing table of Wireguard. A transmission of outbound packet will consult this table to search for the appropriate public key, using the IP destination address of the packet. With inbound packets, after decryption and authentication, the source IP address need to be resolved to the same peer that have the same public key used in the secure session for decrypting the packet. In cryptokey routing table, each peer may optionally pre-specify an outer external IP address and UDP port of that peer's endpoint. If the endpoint is not specified, the external source IP of the machine will be used

to determine the endpoint . This design also allow the peers to roam freely between different external IP addresses as the public is the main identification of the peer.

Combining the roaming and the cryptokey routing table, when the network interface needs to send the data out, the flow start with the inspection of the IP destination of the packet to find the corresponding peer and its public key (An ICMP "no route to host" will be returned in case of no peer). The packet is then encrypted with mentioned AEAD, prepended with extra headers and sent as a UDP packet to the internet. On the receiving flow, when a UDP packet is received, Wireguard finds the matching peer, decrypts the packet and updates the peers' endpoint. If the packet is either not an IP packet or there is no entry for source IP address of the packet inside the cryptokey routing table, the packet is dropped. Otherwise, the packet will be dispatched to the layer above.

4.2 Noise

Noise [Per] is a framework to design secure protocol based on Diffie-Hellman(DH) key exchange [DH76]. With Noise, designers can construct new protocols that offer features such as identity hiding, mutual authentication, and forward secrecy for deriving shared symmetric keys. Furthermore, this framework has gone through formal analysis and verification [KNB18]. Although not all Noise protocols ensure

these security guarantees, the IKpsk2 form used by WireGuard does. The rest of this section gives a general introduction to Noise, without focusing specifically on WireGuard.

A Noise protocol begins with the exchange of **handshake messages** between two parties to generate a shared secret key for the encryption of **transport messages**. The first one to send the handshake message is referred to as the **initiator**, while the recipient of this message is called the **responder**.

Each handshake in the Noise framework involves a long-term **static key pair** and/or **ephemeral key pair**. Every new handshake requires a fresh creation of the ephemeral keys. The static keys may be used for the purpose of authentication, when a peer has little information on the other's identity. Within the handshake, 256-bit **pre-shared symmetric key** can also be mixed in to bind the Noise session to a previous protocol-specific communication. Both the ephemeral and static keys are used in a chain of Diffie-Hellman computation to guarantee the forward secrecy. A cryptographic binding between the final shared secret and the handshake are achieved by the fact that all handshake messages and key material are used to derive the symmetric keys for handshake and the encryption of transport data.

In Noise, the sequential exchange that comprises a handshake message is called a **handshake pattern**. Each handshake pattern consists of **message pattern**, which is arranged from a set of pre-defined **tokens**. The tokens describes the modification to the cipher and hand-

shake state, and data to read or write. The handshake state holds both local key pairs as well as the remote ephemeral and static public keys of peer. The cipher state is made of a key and a counter-based nonce for encryption, a handshake hash and a chaining key to derive the final symmetric key. Noise framework defines the following tokens for the message patterns:

- **e**: The sender creates a new ephemeral key pair, writes the public key into the message buffer, and hashes this key into the handshake hash. In a PSK handshake, the chaining key is mixed with the ephemeral public key.
- **s**: The sender encrypts the static public key and write the result to message buffer. if the handshake symmetric key exists, otherwise the plaintext of the public key will be written. The result will also be hashed into the handshake hash.
- **ee, es, se, ss**: DH computation is performed using the sender's ephemeral or static private key (determined by the first letter) and the receiver's ephemeral or static public key (determined by the second letter). The output is hashed into the chaining key, resetting the cipher state with a new key derived from the chaining key and a zero nonce.
- **psk**: A pre-shared symmetric key is hashed both into the chaining key and the handshake hash state. As the chaining key derives the symmetric key directly, it's essential to mix the nonce into

the chaining key for the generation of a randomized symmetric encryption key. This mitigates the key reuse vulnerability. The ephemeral public key provided by the "e" token serves as this nonce. Consequently, in valid Noise protocols that utilize the **psk** token, encryption must be preceded by the **e** token. An example of invalid pattern is "**psk, s, e**". Here, **s** establishes the encrypted static public key with a fixed symmetric key derived from the PSK.

The order of tokens within a message pattern is crucial. For instance, "**s, ss**" differs from "**ss, s**". In the former, the sender's long-term static public key is presented in plain text ("**s**"), followed by a DH computation between the static keys of both parties ("**ss**"). Any subsequent messages following this pattern will be encrypted with derived keys. In contrast, the latter pattern ("**ss, s**") indicates that the sender's long-term static public key is encrypted, using keys derived from a hash that incorporates the DH computation between the two parties' static keys ("**ss**").

Before the handshake, protocols may already possess information about the peer's static and/or ephemeral public keys. This knowledge is represented by one of the **pre-message patterns**: "**e**," "**s**," "**e, s**," or none if there are no pre-messages. The "**s**" pre-message pattern signifies knowledge of the other party's static public key, while the "**e**" pre-message pattern, used in a fallback scenario, indicates a previously shared fresh ephemeral public key. None of these public keys will be

published again during the handshake; they remain implicit. The pre-message patterns and message patterns together create a **handshake pattern** that includes:

- A pre-message pattern for the initiator, which conveys information about the responder's public keys.
- A pre-message pattern for the responder, which conveys information about the initiator's public keys.
- A chain of message patterns the form the handshake messages.

Each message pattern has a specific direction that influences how the tokens are interpreted. For instance, for the initiator, the $\rightarrow s$ notation signifies sending the initiator's static public key, whereas the responder interprets it as reading the initiator's static public key.

Without knowledge of the exact public keys from the pre-message patterns, it will be impossible to derive the correct symmetric keys during the handshake. Consequently, this will lead to a failure of the handshake due to mismatched handshake states.

Noise specifies several standard handshake patterns, each with different requirements, properties, and security guarantees. Transmission of user data can happen after a completed handshake or inside the buffer of message pattern within a handshake. Nevertheless, sending such data beforehand can lead to reduced authenticity and confidentiality.

A instantiation of a concrete Noise protocol requires a choosing of **handshake pattern**, a **DH function**, a **cipher function**, and a **hash function**. A prologue byte pattern can be established to associate the cryptographic keys with this value. Following is the instantiation of Wireguard:

- **handshake pattern** IKpsk2
- **DH function** X25519
- **cipher function** ChaCha20-Poly1305
- **hash function** BLAKE2s
- **prologue** WireGuard v1 zx2c4 Jason@zx2c4.com

The name of the **IKpsk2** handshake pattern is explained as:

- **I**: static public key is **immediately** transmitted to the responder, despite absent or reduced identity hiding.
- **K**: static public key for the responder is **known** to the initiator.
- **psk2**: a **PSK** is used at the end of the **second** handshake message.

The name for the Noise protocol is concatenated from the string Noise, the handshake pattern, and functions, separated by an underscore (_). Hence, for Wireguard, the full Noise protocol name is: Noise_IKpsk2_25519_ChaChaPoly_BLAKE2s. The details of the IKpsk2 is presented in section 4.2.1.

4.2.1 IKpsk2 Handshake Pattern

The steps to execute the IKpsk2 handshake pattern used in Wireguard will be elaborated in this section. Following is the explanation for several operators, functions and symbols:

- S_i^{pub} : the initiator's 32-byte static public key.
- E_r^{priv} : the responder's 32-byte ephemeral private key.
- Q : optional 32-byte pre-shared symmetric key.
- T_i^{send}, T_r^{recv} : the initiator's 32-byte transport data symmetric key for sending, and the responder's 32-byte transport data symmetric key for receiving.

The symbols for handshake and cipher state:

- h : the 32-byte handshake hash.
- ck : the 32-byte chaining key.
- k : the 32-byte handshake cipher key.

The functions and constants that will be used:

- **ECDH(private key, public key)**: the Elliptic Curve Diffie-Hellman function [Ber06], returning the 32-byte shared secret from 32-byte public key and 32-byte private key.
- **ECDH_GEN()**: returns a key pairs consisting of a random Curve25519 private key and the related public key.

- **AEAD_ENC(cleartext, aad, key, nonce)**: applies ChaCha20-Poly1305 AEAD encryption on a variable-length cleartext message, using a 256-bit key, a 96-bit counter after prepending the 64-bit nonce with 4 bytes of zero and arbitrary length additional authenticated data to return the ciphertext including a 16-byte authentication tag.
- **AEAD_DEC(ciphertext, aad, key, nonce)**: the reverse of **AEAD_ENC** that authenticates and decrypt the ciphertext to return the original plaintext.
- **Hash(data)**: hashes the variable-length data with BLAKE2s function to return 32 bytes of output.
- **KDF_n(key, input)**: derives the n 32-bytes keys with HKDF, using unkeyed BLAKE2s as the hash function for the HMAC [PP09, section 13.2]. The construction of HKDF is as follow:

$$PRK = HMAC(key, input)$$

$$T_1 = HMAC(PRK, 0x1)$$

$$T_2 = HMAC(PRK, T_1 || 0x2)$$

...

$$T_n = HMAC(PRK, T_{n-1} || n)$$

The IKpsk2 handshake pattern in compact notation are shown Figure 4.2. The details of the computation follows below, as a reiteration

in [Wu19, p. 15].

IKpsk2:

1. $\leftarrow s$
...
2. $\rightarrow e, es, s, ss$
3. $\leftarrow e, ee, se, psk$

Figure 4.2: Compact notation for IKpsk2. The dots distinguish between pre-messages and messages. See [Per, p. 36]

The hashing of protocol name and prologue begins the handshake computation.

1. $h_0 = \mathbf{Hash}(\text{"Noise IKpsk2 25519 ChaChaPoly BLAKE2s"})$

2. $ck_0 = h_0$

3. $h_1 = \mathbf{Hash}(h_0 \parallel \text{"WireGuard v1 zx2c4 Jason@zx2c4.com"})$

Next, the handshake messages and pre-messages will be processed. Since there are message tokens that have different interpretation depending on the initiator and the responder, I: and R: will be prefixed before the line to constrain them to the initiator or responder correspondingly. Following are how to interpret the three (pre-)messages from Figure 4.2:

0. $\leftarrow s$: a pre-message from responder to initiator, with “s” as the previously out-of-band exchanged static public key S_r^{pub} .

4. $h_2 = \mathbf{Hash}(h_1 \parallel S_r^{pub})$

1. $\rightarrow e, es, s, ss$: the initiator’s message to the responder:

- “e”: the initiator creates an ephemeral key pair, transmits the public key E_i^{pub} to the responder. E_i^{pub} is mixed into the handshake hash and the chaining key because of the PSK.

5. I: $(E_i^{priv}, E_i^{pub}) = \mathbf{ECDH_GEN}()$; write E_i^{pub}

R: read E_i^{pub}

6. $h_3 = \mathbf{Hash}(h_2 \parallel E_i^{pub})$

7. $ck_1 = \mathbf{KDF}_1(ck_0, E_i^{pub})$

- “es”: the responder calculates the $es = \mathbf{ECDH}(S_r^{priv}, E_i^{pub})$.

The same shared secret is calculated as $es = \mathbf{ECDH}(E_i^{priv}, S_r^{pub})$

for the initiator.

8. I: $(ck_2, k_0) = \mathbf{KDF}_2(ck_1, \mathbf{ECDH}(E_i^{priv}, S_r^{pub}))$

R: $(ck_2, k_0) = \mathbf{KDF}_2(ck_1, \mathbf{ECDH}(S_r^{priv}, E_i^{pub}))$

- “s”: the initiator transmits its encrypted static public key S_i^{pub} . On failure of decryption, the handshake is aborted.

9. I: $msg = \mathbf{AEAD_ENC}(S_i^{pub}, h_3, k_0, 0)$; write msg

R: $S_i^{pub} = \mathbf{AEAD_DEC}(msg, h_3, k_0, 0)$; abort on failure

10. $h_4 = \mathbf{Hash}(h_3 \parallel msg)$

- “ss”: the responder calculates the $ss = \mathbf{ECDH}(S_r^{priv}, S_i^{pub})$.

The same shared secret is calculated as $ss = \mathbf{ECDH}(S_i^{priv}, S_r^{pub})$

for the initiator.

11. I: $(ck_3, k_1) = \mathbf{KDF}_2(ck_2, \mathbf{ECDH}(S_i^{priv}, S_r^{pub}))$

R: $(ck_3, k_1) = \mathbf{KDF}_2(ck_2, \mathbf{ECDH}(S_r^{priv}, S_i^{pub}))$

- Add an encrypted (possibly) empty payload to the message buffer in the handshake. If the decryption fails for the respon-

der, it stops the handshake. For Wireguard, the payload is the current timestamp.

12. I: $c = \mathbf{AEAD_ENC}(\text{time}, h_4, k_1, 0)$; write c

R: $\text{time} = \mathbf{AEAD_DEC}(c, h_4, k_1, 0)$; abort on failure

13. $h_5 = \mathbf{Hash}(h_4 \parallel m)$

2. $\leftarrow e, ee, se, \text{psk}$: the responder's message to the initiator:

- “e”: the responder creates an ephemeral key pair, transmits the public key E_r^{pub} to the initiator. E_r^{pub} is mixed into the handshake hash and the chaining key because of the PSK.

14. R: $(E_r^{priv}, E_i^{pub}) = \mathbf{ECDH_GEN}()$; write E_r^{pub}

I: read E_r^{pub}

15. $h_6 = \mathbf{Hash}(h_5 \parallel E_r^{pub})$

16. $ck_4 = \mathbf{KDF}_1(ck_3, E_r^{pub})$

- “ee”: the responder calculates the $ee = \mathbf{ECDH}(E_r^{priv}, E_i^{pub})$.

The same secret is calculated as $es = \mathbf{ECDH}(E_i^{priv}, E_r^{pub})$ for the initiator.

17. R: $ck_5 = \mathbf{KDF}_1(ck_4, \mathbf{ECDH}(E_r^{priv}, E_i^{pub}))$

I: $ck_5 = \mathbf{KDF}_1(ck_4, \mathbf{ECDH}(E_i^{priv}, E_r^{pub}))$

- “se”: the responder calculates the $se = \mathbf{ECDH}(E_r^{priv}, S_i^{pub})$.

The same secret is calculated as $se = \mathbf{ECDH}(S_i^{priv}, E_r^{pub})$ for the initiator.

18. R: $ck_6 = \mathbf{KDF}_1(ck_5, \mathbf{ECDH}(E_r^{priv}, S_i^{pub}))$

I: $ck_6 = \mathbf{KDF}_1(ck_5, \mathbf{ECDH}(S_i^{priv}, E_r^{pub}))$

- “psk”: include the pre-shared symmetric key Q in both chaining key and the handshake hash.

$$19. (ck_7, \tau, k_2) = \mathbf{KDF}_3(ck_6, Q)$$

$$20. h_7 = \mathbf{Hash}(h_6 \parallel \tau)$$

- Add an encrypted (possibly) empty payload to the message buffer in the handshake. If the decryption fails for the responder, it aborts the handshake.

$$21. \text{R: } c = \mathbf{AEAD_ENC}(m, h_7, k_2, 0); \text{ write } c$$

$$\text{I: } m = \mathbf{AEAD_DEC}(c, h_7, k_2, 0); \text{ abort on failure}$$

$$22. h_8 = \mathbf{Hash}(h_7 \parallel m)$$

At last, the handshake ends with the derivation of the symmetric keys for transport data encryption, using HKDF with the chaining key and a empty string ϵ .

$$23. \text{I: } (T_i^{send}, T_i^{recv}) = \mathbf{KDF}_2(ck_2, \epsilon)$$

$$\text{R: } (T_r^{send}, T_r^{recv}) = \mathbf{KDF}_2(ck_2, \epsilon)$$

This wraps up the execution and overview of the Noise IKpsk2 protocol. Next section will delve into how Wireguard uses this pattern for its payload.

4.3 Wireguard Messages

This sections details the underlying Wireguard messages that realizes the Noise handshake protocol, protects again denial of service attacks and encapsulates the IP packet in a robust tunnel. In Wireguard, there are four kinds of message, all of which start with a single type byte, followed by three zero reserved bytes. This layout allows the

implementation to read the first 4 bytes of the message as a little-endian integer.

type := 0x1 (1 byte)	reserved := 0 ³ (3 bytes)
sender := I_i (4 bytes)	
ephemeral (32 bytes)	
static ($\widehat{32}$ bytes)	
timestamp ($\widehat{12}$ bytes)	
mac1 (16 bytes)	mac2 (16 bytes)

(a) Handshake Initiation

type := 0x2 (1 byte)	reserved := 0 ³ (3 bytes)
sender := I_r (4 bytes)	receiver := I_i (4 bytes)
ephemeral (32 bytes)	
empty ($\widehat{0}$ bytes)	
mac1 (16 bytes)	mac2 (16 bytes)

(b) Handshake Response

Figure 4.3: Handshake messages, see [Don20, p. 10].

Figure 4.3 show first 2 types of Wireguard messages: **Handshake initiation** and **Handshake response**. These 2 messages are the adaptation of Wireguard for Noise handshake. Both of messages contains a **sender index** to indicate a 32-bit index that locally represents the sending peer. The **receiver index** of handshake response is the received **sender index** from the initiator. The ephemeral fields are the ephemeral public keys for key exchange, while the static field of the **handshake initiation** is the encrypted static public key of the initiator after the DH computation. The initiation message also includes an encrypted TAI64N [Ber] timestamp to protect against replay attack. The motivation for the final MAC fields are explained in section 5.4.

These MAC fields are computed using the keyed BLAKE2s MAC with a 16-byte hash value as output. Depending on the recipient and field type, the input and key for each MAC field will be different. To be specific:

MAC1 Using $S_{m'}^{pub}$ as the receiver static public key, the key of the MAC computation is the value: **Hash**("mac1----" || $S_{m'}^{pub}$). The input of the MAC is all the bytes of the message prior to MAC1 field.

MAC2 Given that the latest cookie was received within 120 seconds, this cookie would be the MAC2 key. Otherwise, if the key is too old or there is no such cookie, the MAC2 field will be zeroed out. The data for MAC2 is all the bytes prior to the MAC2 field.

The third type of message is the **Cookie Reply Message**, which is sent when the current peer is under load. The exact condition to transmit this cookie message is after the reception of a handshake message with valid MAC1 but invalid/expired MAC2 and the peer is under load. The format of the message is as follow:

type := 0x3 (1 byte)	reserved := 0 ³ (3 bytes)
receiver := $I_{m'}$ (4 bytes)	
nonce := ρ^{24} (24 bytes)	
cookie ($\hat{16}$ bytes)	

Figure 4.4: Cookie Reply Message, see [Don20, p. 13].

The receiver is the sender from the previous handshake message.

With a random secret value R_m that changes every two minutes, the concatenation of the last message's external IP source address and its UDP source port, and MAC1 as M, the encrypted cookie is created like below:

$$\tau = \mathbf{BLAKE2s}(R_m, A_{m'})$$

$$cookie = \mathbf{XAEAD}(\tau, M, \mathbf{Hash}("cookie--" \| S_m^{pub}), nonce)$$

By using M as the additional authenticated data field, the cookie reply is securely bound to the corresponding message, preventing peers from being targeted by fraudulent cookie reply attacks. Additionally, this message is smaller than both the handshake initiation and handshake response messages, reducing the risk of amplification attacks.

The final type of messages is the Transport Data Message - the encrypted encapsulated payload for secure communication. A counter is included in the packet to use as the nonce for AEAD and a tool to prevent replay attack. The inner IP packet will padded with zero bytes until its length is a multiple of 16 to achieve a better address alignment, thus improving the performance on many CPU architectures.

type := 0x4 (1 byte)	reserved := 0 ³ (3 bytes)
receiver := $I_{m'}$ (4 bytes)	
counter (8 bytes)	
packet ($\widehat{\ P\ }$ bytes)	

Figure 4.5: Transport Data Message, see [Don20, p. 12].

4.4 Timers

Chapter 5

Wireguard and the IoT

5.1 Cryptography Resilience

5.2 Virtuous Silence

5.3 Stateful firewall

5.4 Cookies and Denial of Services Attack

5.5 Roaming & NDP

5.6 Message Lengths & Fragmentation

Chapter 6

Related Work

Chapter 7

Design and Implementation of Wireguard for GNRC

Chapter 8

Testing

8.1 Methodology

8.2 Environment

8.3 Scenario 1

8.4 Scenario 2

8.5 Scenario 3

8.6 Scenario 4

8.7 Memory Usage

Chapter 9

Conclusion and Future Work

References

- [Arc20] Scott Arciszewski. *XChaCha: eXtended-nonce ChaCha and AEAD_XChaCha20_Poly1305*. Internet-Draft draft-irtf-cfrg-xchacha-03. Work in Progress. Internet Engineering Task Force, Jan. 2020. 18 pp. URL: <https://datatracker.ietf.org/doc/draft-irtf-cfrg-xchacha/03/>.
- [Bac+18] Emmanuel Baccelli et al. “RIOT: An Open Source Operating System for Low-End Embedded Devices in the IoT”. In: *IEEE Internet of Things Journal* 5.6 (2018), pp. 4428–4440. DOI: 10.1109/JIOT.2018.2815038.
- [Ber] Daniel J. Bernstein. *TAI64, TAI64N, and TAI64NA*. URL: <https://cr.yp.to/libtai/tai64.html>.
- [Ber05] Daniel J. Bernstein. *The Poly1305-AES message-authentication code*. 2005. URL: <https://cr.yp.to/mac/poly1305-20050329.pdf>.
- [Ber06] Daniel J. Bernstein. “Curve25519: New Diffie-Hellman Speed Records”. In: *Public Key Cryptography - PKC 2006*. Ed. by Moti Yung et al. Berlin, Heidelberg: Springer Berlin Heidelberg, 2006, pp. 207–228. ISBN: 978-3-540-33852-9.
- [Ber08] Daniel J. Bernstein. *ChaCha, a variant of Salsa20*. 2008. URL: <https://cr.yp.to/chacha/chacha-20080120.pdf>.
- [Bor+12] Carsten Bormann et al. *Neighbor Discovery Optimization for IPv6 over Low-Power Wireless Personal Area Networks (6LoWPANs)*. RFC 6775. Nov. 2012. DOI: 10.17487/RFC6775. URL: <https://www.rfc-editor.org/info/rfc6775>.

- [Bru16] Simon Brummer. “Concept and Implementation of TCP for the RIOT Operating System and Evaluation of Common TCP-Extensions for the Internet of Things”. In: (2016).
- [DH17] Dr. Steve E. Deering and Bob Hinden. *Internet Protocol, Version 6 (IPv6) Specification*. RFC 8200. July 2017. DOI: 10.17487/RFC8200. URL: <https://www.rfc-editor.org/info/rfc8200>.
- [DH76] W. Diffie and M. Hellman. “New directions in cryptography”. In: *IEEE Transactions on Information Theory* 22.6 (1976), pp. 644–654. DOI: 10.1109/TIT.1976.1055638.
- [Don20] Jason A. Donenfeld. *WireGuard: Next Generation Kernel Network Tunnel*. 2020. URL: <https://wireguard.com/papers/wireguard.pdf>.
- [GC06] Mukesh Gupta and Alex Conta. *Internet Control Message Protocol (ICMPv6) for the Internet Protocol Version 6 (IPv6) Specification*. RFC 4443. Mar. 2006. DOI: 10.17487/RFC4443. URL: <https://www.rfc-editor.org/info/rfc4443>.
- [GCS21] Carles Gomez, Jon Crowcroft, and Michael Scharf. *TCP Usage Guidance in the Internet of Things (IoT)*. RFC 9006. Mar. 2021. DOI: 10.17487/RFC9006. URL: <https://www.rfc-editor.org/info/rfc9006>.
- [Hen11] Sushil Jajodia Henk C. A. van Tilborg, ed. *Encyclopedia of Cryptography and Security*. 2nd ed. Springer New York, NY, 2011. DOI: <https://doi.org/10.1007/978-1-4419-5906-5>.
- [HH05] Brian Haberman and Bob Hinden. *Unique Local IPv6 Unicast Addresses*. RFC 4193. Oct. 2005. DOI: 10.17487/RFC4193. URL: <https://www.rfc-editor.org/info/rfc4193>.
- [KE10] Hugo Krawczyk and Pasi Eronen. *HMAC-based Extract-and-Expand Key Derivation Function (HKDF)*. RFC 5869. May 2010. DOI: 10.17487/RFC5869. URL: <https://www.rfc-editor.org/info/rfc5869>.

- [KNB18] Nadim Kobeissi, Georgio Nicolas, and Karthikeyan Bhargavan. *Noise Explorer: Fully Automated Modeling and Verification for Arbitrary Noise Protocols*. Cryptology ePrint Archive, Paper 2018/766. 2018. URL: <https://eprint.iacr.org/2018/766>.
- [Len16] M. Lenders. “Analysis and Comparison of Embedded Network Stacks”. MA thesis. Freien Universität Berlin, 2016.
- [Mon+07] Gabriel Montenegro et al. *Transmission of IPv6 Packets over IEEE 802.15.4 Networks*. RFC 4944. Sept. 2007. DOI: 10.17487/RFC4944. URL: <https://www.rfc-editor.org/info/rfc4944>.
- [NJT07] Dr. Thomas Narten, Tatsuya Jinmei, and Dr. Susan Thomson. *IPv6 Stateless Address Autoconfiguration*. RFC 4862. Sept. 2007. DOI: 10.17487/RFC4862. URL: <https://www.rfc-editor.org/info/rfc4862>.
- [NL18] Yoav Nir and Adam Langley. *ChaCha20 and Poly1305 for IETF Protocols*. RFC 8439. June 2018. DOI: 10.17487/RFC8439. URL: <https://www.rfc-editor.org/info/rfc8439>.
- [Per] Trevor Perrin. *The Noise Protocol Framework*. revision 34, 2018-07-11. URL: <https://noiseprotocol.org/noise.pdf>.
- [Pet+15] Hauke Petersen et al. *Old Wine in New Skins? Revisiting the Software Architecture for IP Network Stacks on Constrained IoT Devices*. 2015. arXiv: 1502.01968 [cs.NI]. URL: <https://arxiv.org/abs/1502.01968>.
- [Pos80] J. Postel. *User Datagram Protocol*. RFC 768. 1980. DOI: 10.17487/RFC0768. URL: <https://www.rfc-editor.org/info/rfc768>.
- [PP09] Christof Paar and Jan Pelzl. *Understanding Cryptography: A Textbook for Students and Practitioners*. 1st. Springer Publishing Company, Incorporated, 2009. ISBN: 3642041000.

- [SA15] Markku-Juhani O. Saarinen and Jean-Philippe Aumasson. *The BLAKE2 Cryptographic Hash and Message Authentication Code (MAC)*. RFC 7693. Nov. 2015. DOI: 10.17487/RFC7693. URL: <https://www.rfc-editor.org/info/rfc7693>.
- [SB09] Zack Shelby and Carsten Bormann. *6LoWPAN The Wireless Embedded Internet*. WILEY, 2009.
- [Shu+07] Feng Shu et al. “Packet loss analysis of the IEEE 802.15.4 MAC without acknowledgements”. In: *IEEE Communications Letters* 11.1 (2007), pp. 79–81. DOI: 10.1109/LCOMM.2007.061295.
- [Wu19] Peter Wu. “Analysis of the WireGuard protocol”. MA thesis. Eindhoven University of Technology, 2019.